

Distributed Logical Volume with Striped

环境：本文仅讨论实现，并不涉及性能和安全

IP-SAN(iscsi targets) four server


One clients

平台：Vmware 6.0 ACE

CentOS5 update 2 x86

我的 4 个节点分别为 gfs1 gfs2 gfs3 gfs4

网络拓扑图：



Distributed Logical Volume with Striped

一、前言

首先我们做一个试验：

```
[root@gfs2 ~]# mkfs.ext3 /dev/sdb
```

mke2fs 1.39 (29-May-2006)

/dev/sdb is entire device, not just one partition!

Proceed anyway? (y,n) y

Filesystem label=

OS type: Linux

Block size=4096 (log=2)

Fragment size=4096 (log=2)

1310720 inodes, 2621440 blocks

131072 blocks (5.00%) reserved for the super user

First data block=0

Maximum filesystem blocks=2684354560

80 block groups

32768 blocks per group, 32768 fragments per group

16384 inodes per group

Superblock backups stored on blocks:

32768, 98304, 163840, 229376, 294912, 819200, 884736, 1605632

Writing inode tables: done

Creating journal (32768 blocks): done

Writing superblocks and filesystem accounting information: done

This filesystem will be automatically checked every 20 mounts or 180 days, whichever comes first. Use tune2fs -c or -i to override.

[root@gfs2 ~]# mkdir /test

[root@gfs2 ~]# mount /dev/sdb /test

[root@gfs2 ~]# df -h

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/mapper/VolGroup00-LogVol00	7.0G	2.1G	4.5G	32%	/
/dev/sda1	99M	12M	83M	13%	/boot
tmpfs	189M	0	189M	0%	/dev/shm
/dev/hdc	3.8G	3.8G	0	100%	/media/CentOS_5.2_Final
/dev/sdb	9.9G	151M	9.2G	2%	/test

做这个试验的目的是要大家知道，一块磁碟不一定要分区。一样能使用。但是在测试中 linux 的软 RAID 一定要分区。现在回到主机 gfs1 上去做 LVM 实验。

二、建立一个单级条带的逻辑卷

1、建立 Physical Volumes

```
[root@gfs1 ~]# pvcreate /dev/sdb
```

```
Physical volume "/dev/sdb" successfully created
```

```
[root@gfs1 ~]# pvcreate /dev/sdc
```

```
Physical volume "/dev/sdc" successfully created
```

2、建立 Volume Group

```
[root@gfs1 ~]# vgcreate VG0 /dev/sdb /dev/sdc
```

```
Volume group "VG0" successfully created
```

3、建立一个单级条带的逻辑卷

```
[root@gfs1 ~]# lvcreate -L 2G -n lv0 VG0
```

```
Logical volume "lv0" created
```

4、用 gfs 格式化逻辑卷

```
[root@gfs1 ~]# [root@gfs1 ~]# gfs_mkfs -p lock_nolock -j 1 /dev/VG0/lv0
```

```
This will destroy any data on /dev/VG0/lv0.
```

```
Are you sure you want to proceed? [y/n] y
```

```
Blocksize: 4096
```

```
-bash: This: command not found
```

```
Filesystem Size: 491460
```

```
Journals: 1
```

```
Locking Protocol: lock_nolock
```

```
Lock Table:
```

```
Syncing...
```

```
All Done
```

1.5 挂载格式化完成的文件系统

```
[root@gfs1 ~]# mkdir /gfs_nolock
```

```
[root@gfs1 ~]# mount -t gfs /dev/VG0/lv0 /gfs_nolock/
```

```
[root@gfs1 ~]# df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/mapper/VolGroup00-LogVol00	7.0G	2.1G	4.5G	32%	/
/dev/sda1	99M	12M	83M	13%	/boot
tmpfs	189M	0	189M	0%	/dev/shm
/dev/mapper/VG0-lv0	1.9G	20K	1.9G	1%	/gfs_nolock

```
[root@gfs1 ~]# mount -l -t gfs
```

/dev/mapper/VG0-lv0	on	/gfs_nolock	type	gfs
(rw,localflocks,localcaching,oopses_ok)				

三、建立一个多级条带的逻辑卷

1、建立 Physical Volumes

```
[root@gfs3 ~]# pvcreate /dev/sd[b,c]
```

Physical volume "/dev/sdb" successfully created

Physical volume "/dev/sdc" successfully created

2、建立 Volume Group

```
[root@gfs3 ~]# vgcreate vg0 /dev/sdb /dev/sdc
```

Volume group "vg0" successfully created

3、建立一个多级条带的逻辑卷

```
[root@gfs3 ~]# lvcreate -i2 -l4 -L3G -nlv0 vg0
```

Logical volume "lv0" created

4、用 gfs 格式化逻辑卷

```
[root@gfs3 ~]# gfs_mkfs -plock_nolock -j 1 /dev/vg0/lv0
```

This will destroy any data on /dev/vg0/lv0.

Are you sure you want to proceed? [y/n] y

Device: /dev/vg0/lv0

Blocksize: 4096

Filesystem Size: 753580
Journals: 1
Resource Groups: 12
Locking Protocol: lock_nolock
Lock Table:
Syncing...
All Done

4、挂载格式化完成的文件系统

```
[root@gfs3 ~]# mkdir /testlv
```

```
[root@gfs3 ~]# mount -t gfs /dev/vg0/lv0 /testlv/
```

```
[root@gfs3 ~]# df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/mapper/VolGroup00-LogVol00	7.0G	2.1G	4.5G	32%	/
/dev/sda1	99M	12M	83M	13%	/boot
tmpfs	189M	0	189M	0%	/dev/shm
/dev/hdc	3.8G	3.8G	0	100%	/media/CentOS_5.2_Final
/dev/mapper/vg0-lv0	2.9G	20K	2.9G	1%	/testlv

```
[root@gfs3 ~]# mount -l -t gfs  
/dev/mapper/vg0-lv0 on /testlv type gfs (rw,localflocks,localcaching,oopses_ok)
```

其他的一些操作我就不再多讲了。网上很多大家可以搜索。现在讲下和 IP-SAN 的具体的分布式的应用。大家都知道 LUSTRE 是个分布式的集群文件系统。其实 GFS 本身不是完全分布式的。（这里过气的 GBIND 就不再讨论）他仅仅是一个有 LOCK 机制和多 journal 的文件系统。靠分布式的是它下层的 LVM。大家可能看过红帽的 RHCS 的 OVERVIEW 里面有个讲 LVM2 的 CLUSTER 的图片。相信大家都熟悉 clvmd 这个程式。这个程式运行在 GFS 的 node 上。这个程式的作用仅仅是能让 GFS node 识别 share storage 上的逻辑卷。其实和分布式没有任何一点关系。下面我就尝试一下用 ISCSI+LVM 来逻辑分布存储，如果用单

级条带分布的话，其实没有任何意义。LV 的 I/O 也上不去。

四、Distributed Logical Volume with Striped

1、在4个节点上先把本地磁盘target出来

```
[root@gfs1 ~]# yum install scsi-target-utils
```

```
[root@gfs1 ~]# chkconfig tgtd on
```

```
[root@gfs1 ~]# service tgtd restart
```

Stopping SCSI target daemon:

Starting SCSI target daemon: [OK]

定义 target 的 qualified 的名字

```
[root@gfs1 ~]# tgtadm --lld iscsi --op new --mode target --tid 1 -T  
iqn.2008-12.sys.sdb
```

```
[root@gfs1 ~]# tgtadm --lld iscsi --op new --mode target --tid 2 -T  
iqn.2008-12.sys.sdc
```

为创建目标增加分区

```
[root@gfs1 ~]# tgtadm --lld iscsi --op new --mode logicalunit --tid 1 --lun 1 -b  
/dev/sdb
```

```
[root@gfs1 ~]# tgtadm --lld iscsi --op new --mode logicalunit --tid 2 --lun 1 -b  
/dev/sdc
```

定义客户端的访问

```
[root@gfs1 ~]# tgtadm --lld iscsi --op bind --mode target --tid 1 -I ALL
```

```
[root@gfs1 ~]# tgtadm --lld iscsi --op bind --mode target --tid 2 -I ALL
```

验证

```
[root@gfs1 ~]# tgtadm --lld iscsi --op show --mode target |grep Target
```

Target 1: iqn.2008-12.sys.sdb

Target 2: iqn.2008-12.sys.sdc

到这里我的其他 4 台机器都一样。所以我搞个脚本去运行就可以了。我这里是为了图简便。希望如果你要有什么价值的应用的的话。自己理顺一下每个节点的 target qualified 的名字。

2、调整 client 端，发现：

```
[root@client ~]# iscsiadm -m discovery -t sendtargets -p gfs1
```

```
172.18.174.1:3260,1 iqn.2008-12.sys.sdb
```

```
172.18.174.1:3260,1 iqn.2008-12.sys.sdc
```

```
[root@client ~]# iscsiadm -m discovery -t sendtargets -p gfs2
```

```
172.18.174.2:3260,1 iqn.2008-12.sys.sdb
```

```
172.18.174.2:3260,1 iqn.2008-12.sys.sdc
```

```
[root@client ~]# iscsiadm -m discovery -t sendtargets -p gfs3
```

```
172.18.174.3:3260,1 iqn.2008-12.sys.sdb
```

```
172.18.174.3:3260,1 iqn.2008-12.sys.sdc
```

```
[root@client ~]# iscsiadm -m discovery -t sendtargets -p gfs4
```

```
172.18.174.4:3260,1 iqn.2008-12.sys.sdb
```

```
172.18.174.4:3260,1 iqn.2008-12.sys.sdc
```

```
[root@client ~]# service iscsi restart
```

完成后你就可以去数盘了。哈哈我的是：sd[b,c,d,e,f,g,h,i]八个。

```
Disk /dev/sdb: 10.7 GB, 10737418240 bytes
```

```
Disk /dev/sdc: 17.1 GB, 17179869184 bytes
```

```
Disk /dev/sdd: 10.7 GB, 10737418240 bytes
```

```
Disk /dev/sdf: 10.7 GB, 10737418240 bytes
```

```
Disk /dev/sdg: 17.1 GB, 17179869184 bytes
```

```
Disk /dev/sdh: 17.1 GB, 17179869184 bytes
```

```
Disk /dev/sde: 17.1 GB, 17179869184 bytes
```

```
Disk /dev/sdi: 10.7 GB, 10737418240 bytes
```

这里可以看下 initiator 这端的标示是不规则的。都是靠 UDEV 来扫描生成盘符。

所以你可以调整 UDEV 让其固定盘符。我这里测试就不固定了。最后我分两个 VG 。容量一样的分到一个 VG（为什么，自己做下实验不一样的分下就知道了）

```
[root@client ~]# pvcreate /dev/sd{b,c,d,e,f,g,g,i} -ff
```

```
/dev/cdrom: open failed: Read-only file system
```

```
Attempt to close device '/dev/cdrom' which is not open.
```

```
Physical volume "/dev/sdb" successfully created
```

Physical volume "/dev/sdc" successfully created
Physical volume "/dev/sdd" successfully created
Physical volume "/dev/sde" successfully created
Physical volume "/dev/sdf" successfully created
Physical volume "/dev/sdg" successfully created
Physical volume "/dev/sdg" successfully created
Really INITIALIZE physical volume "/dev/sdi" of volume group "vg0" [y/n]? y
WARNING: Forcing physical volume creation on /dev/sdi of volume group "vg0"
Physical volume "/dev/sdi" successfully created
[root@client ~]# vgcreate iscsi_vg_10g /dev/sd{b,d,f,i}
Volume group "iscsi_vg_10g" successfully created
[root@client ~]# vgcreate iscsi_vg_17g /dev/sd{c,g,h,e}
Volume group "iscsi_vg_17g" successfully created
[root@client ~]# lvcreate -i4 -l4 -l10236 -n10g_lv iscsi_vg_10g
/dev/cdrom: open failed: Read-only file system
Logical volume "10g_lv" created

[root@client ~]# lvcreate -i4 -l4 -l16380 -n17g_lv iscsi_vg_17g

/dev/cdrom: open failed: Read-only file system
Logical volume "17g_lv" created

格式化:

[root@client ~]# mkfs.ext3 /dev/iscsi_vg_10g/10g_lv
mke2fs 1.40.11 (17-June-2008)
Filesystem label=
OS type: Linux
Block size=4096 (log=2)
Fragment size=4096 (log=2)
5242880 inodes, 10481664 blocks

524083 blocks (5.00%) reserved for the super user

First data block=0

Maximum filesystem blocks=0

320 block groups

32768 blocks per group, 32768 fragments per group

16384 inodes per group

Superblock backups stored on blocks:

32768, 98304, 163840, 229376, 294912, 819200, 884736, 1605632,
2654208,
4096000, 7962624

Writing inode tables: done

Creating journal (32768 blocks): done

Writing superblocks and filesystem accounting information: done

This filesystem will be automatically checked every 28 mounts or
180 days, whichever comes first. Use tune2fs -c or -i to override.

五、测试

我这里格式化一个做测试就够了。

```
[root@client /]# mkdir lvm
```

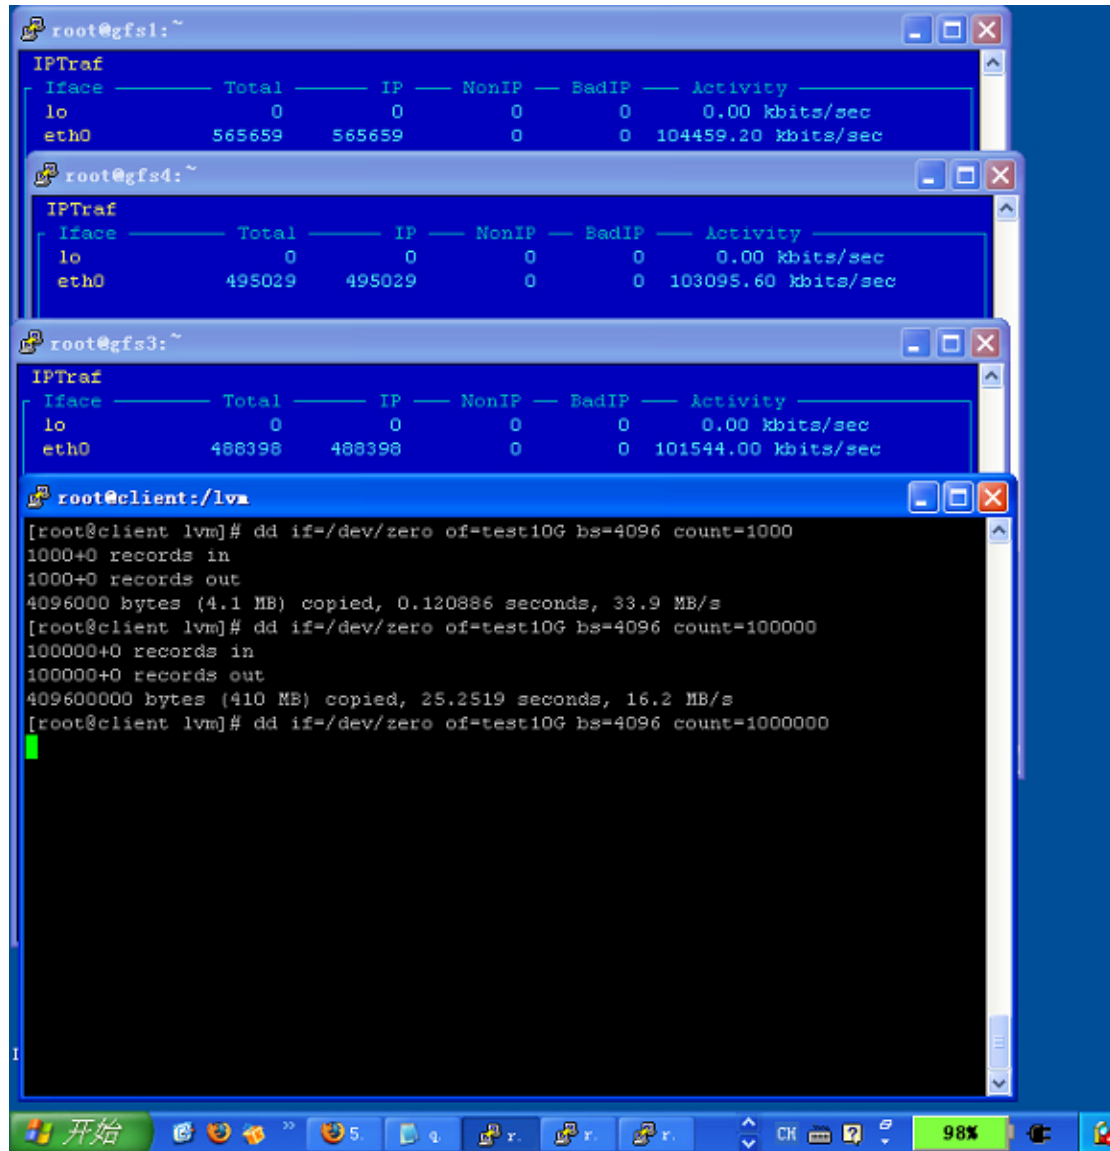
```
[root@client /]# mount /dev/iscsi_vg_10g/10g_lv /lvm
```

```
[root@client /]# df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/mapper/VolGroup00-LogVol00	7.2G	2.2G	4.7G	32%	/
/dev/sda1	99M	18M	77M	19%	/boot
tmpfs	62M	0	62M	0%	/dev/shm
/dev/mapper/iscsi_vg_10g-10g_lv					

40G 177M 38G 1% /lvm

4 个 ISCSI TARGETS 的节点都开启终端。用 iptarf 监控流量。由于屏幕的原因我只监视了 3 个节点。（最后再申明我只是一个实验，没有涉及性能和安全。）



The screenshot displays four terminal windows stacked vertically. The top three windows show the output of the 'IPTraf' command, which provides a summary of network traffic for the 'lo' and 'eth0' interfaces. The bottom window shows the execution of 'dd' commands to copy data from '/dev/zero' to a file named 'test10G'.

Terminal 1 (root@gfs1:~):

```
IPTraf
-----
Iface  Total      IP  NonIP  BadIP  Activity
-----
lo      0          0      0      0      0.00 kbits/sec
eth0    565659    565659  0      0    104459.20 kbits/sec
```

Terminal 2 (root@gfs4:~):

```
IPTraf
-----
Iface  Total      IP  NonIP  BadIP  Activity
-----
lo      0          0      0      0      0.00 kbits/sec
eth0    495029    495029  0      0    103095.60 kbits/sec
```

Terminal 3 (root@gfs3:~):

```
IPTraf
-----
Iface  Total      IP  NonIP  BadIP  Activity
-----
lo      0          0      0      0      0.00 kbits/sec
eth0    488398    488398  0      0    101544.00 kbits/sec
```

Terminal 4 (root@client:/lvm):

```
[root@client lvm]# dd if=/dev/zero of=test10G bs=4096 count=1000
1000+0 records in
1000+0 records out
4096000 bytes (4.1 MB) copied, 0.120886 seconds, 33.9 MB/s
[root@client lvm]# dd if=/dev/zero of=test10G bs=4096 count=100000
100000+0 records in
100000+0 records out
409600000 bytes (410 MB) copied, 25.2519 seconds, 16.2 MB/s
[root@client lvm]# dd if=/dev/zero of=test10G bs=4096 count=1000000
```

(qq174375@gmail.com)